



MedSAM Segment Anything in Medical Images

A source-faithful paper talk from Ma et al. 2024

Research Presentation Program

Department of Statistics & Data Science

2026-08-29



香港中文大學
The Chinese University of Hong Kong

The evaluation object is overlap under a prompted mask

$$\text{DSC}(G, S) = \frac{2|G \cap S|}{|G| + |S|}$$

Here G is the expert annotation and S is the predicted segmentation mask; the paper reports DSC for internal and external validation tasks.

Model components


- ground-truth mask G
- segmentation mask S
- Dice similarity coefficient

Interpretation

MedSAM is judged by segmentation overlap after a bounding-box prompt specifies the target region.

Source: Ma et al. 2024 Methods, DOI 10.1038/s41467-024-44824-z.

Internal validation pairs task distributions with real qualitative cases



The source figure shows both task-level DSC/rank behavior and real article medical-image examples.

Figure 3, Ma et al. 2024, CC BY 4.0.

MedSAM keeps SAM's prompt path but trains it on medical masks




Figure 2. Modality distribution to the promptable SAM-style architecture used for medical segmentation.

The paper's claim is broad utility, with validation boundaries

Failure evidence

CT, MRI, and endoscopy dominate the training set, so modality imbalance remains a source-discussed limitation.




Training diversity

motivates

Manipulate sampling

- inspect unseen targets
- preserve prompt semantics
- validate by modality

Comparator arms



Decision rule

go/no-go

Treat MedSAM as a broad promptable foundation model whose deployment claim still depends on target-specific validation.

The deck ends with the article's balance: broad performance evidence plus explicit validation boundaries.

The external CT case is a same-patient segmentation comparison

