


Endpoint Choice Changes the Method Ranking

The method story changes when the endpoint changes.



Interpretation
Calibrated wins recall; burden error is lower-is-better and lowest for Calibrated.
Baseline wins Dice only.


Meeting decision
Freeze endpoint priority before ranking methods.

Average Dice Can Hide the Hard Case


02/04

A case page binds image, GT, prediction, error overlay, and metric.


Synthetic image




GT mask



Prediction



FP/FN overlay



Dice=0.77 lesion recall=0.76 burden error=-2.6%


TP pixels=1372, FP=376, FN=422; failure mechanism: small false-positive island plus shifted lesion boundary.

Scientific object: same synthetic case, aligned GT/prediction/error overlay, case metric next to visual.

Reference retrieval: failure-case query selected inspected pages RRL-022, RRL-024, RRL-018, RRL-013, RRL-025; trace in EVIDENCE_MANIFEST; style not copied.

Limited-Information Multi-Center Experiment

Raw data stay local; shared summaries are judged by endpoint-level evidence.




Comparator path: local-only uses the same local scores without summary sharing; both branches enter the same endpoint gate.

AI Phenotype Error Changes Inference

Segmentation output is a measured variable whose error can attenuate inference.

Target estimand



Toy calibration: $E[T | M, center] = \alpha_{center} + \gamma_{center} M$. Report beta1 after validation-calibrated correction, not a raw AI burden coefficient.

Evidence boundary
This page is a generated-variable inference mechanism,
not completed clinical validity evidence.